Two-solver diagnostic: accuracy and output cost

64 shared question groups; 128 paired states per solver; two draws per mode.

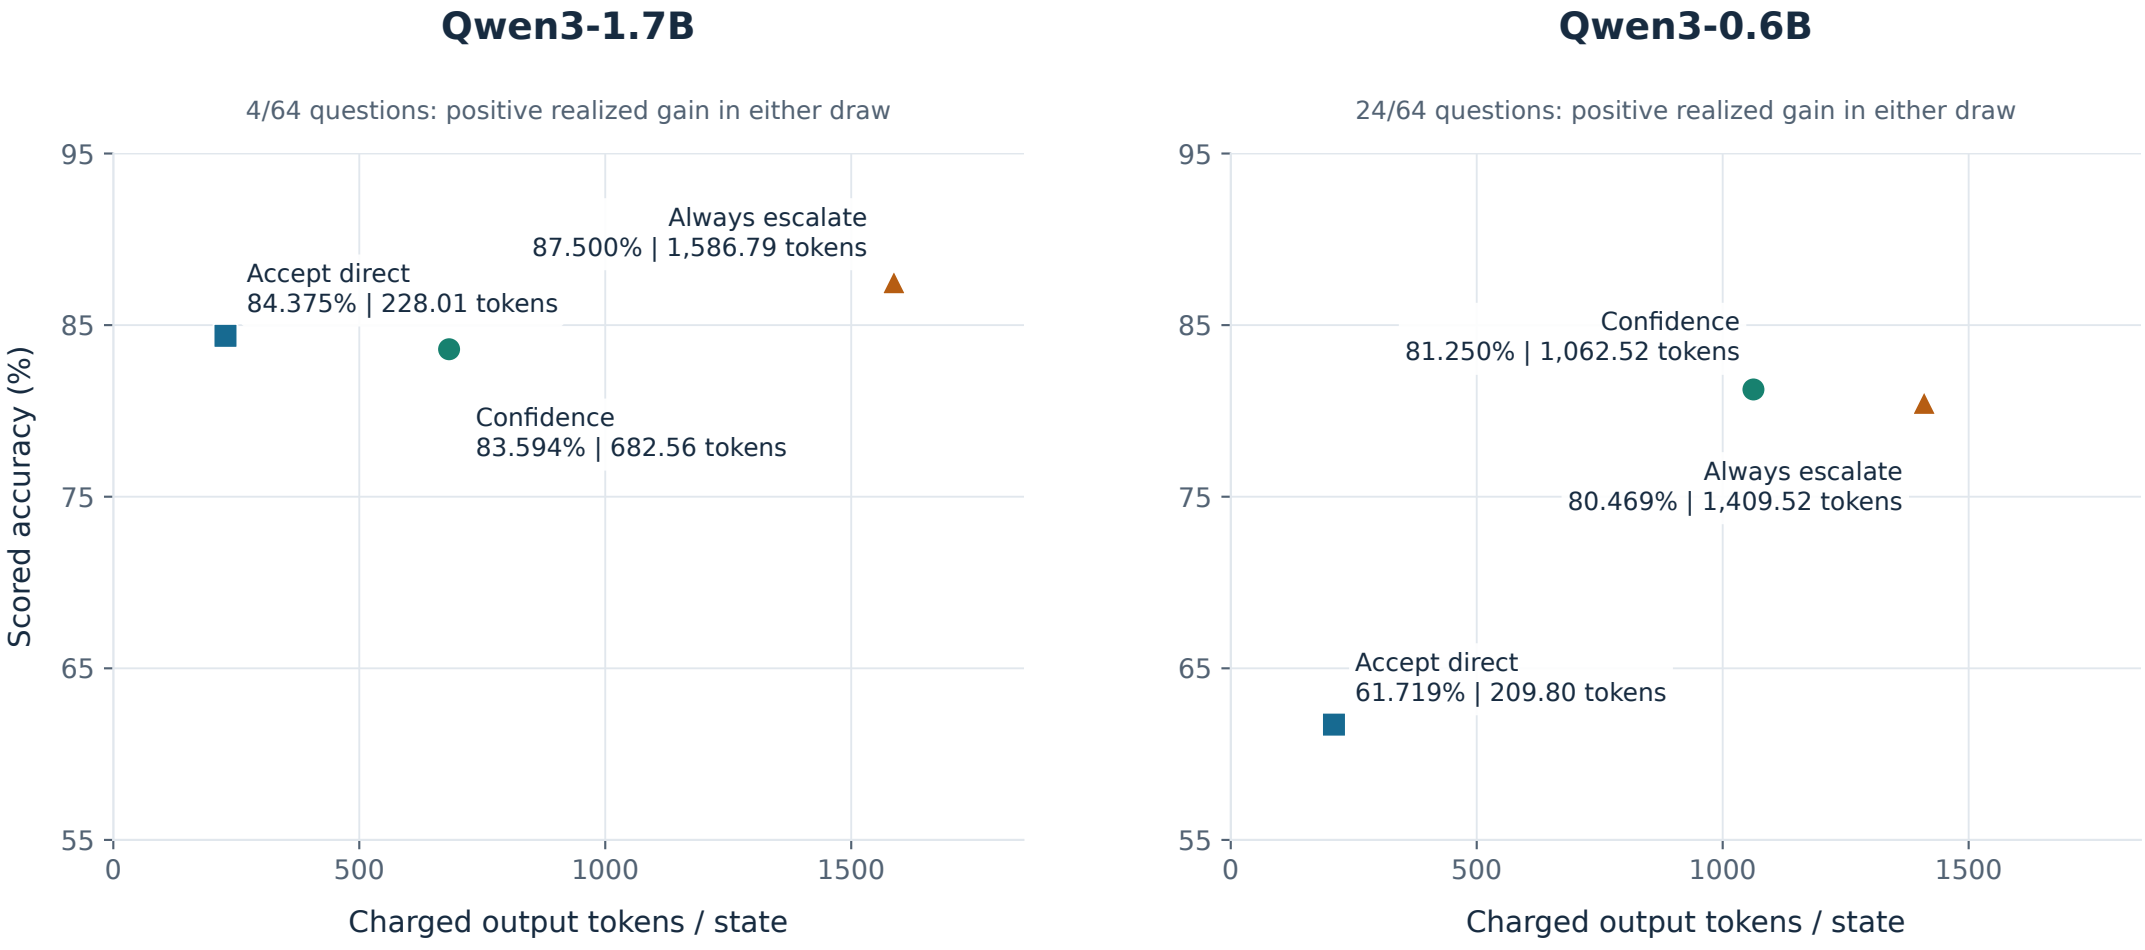

Each label is a reported point estimate. No accuracy/token uncertainty bars are inferred.

Escalation includes the paid direct draft. Tokens are not equal compute across solver sizes.