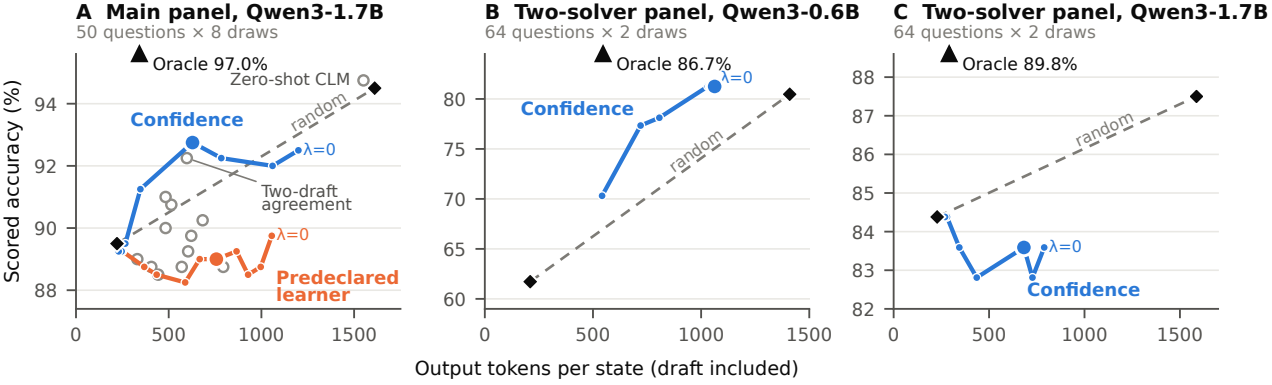
A Main panel, Qwen3-1.7B
50 questions × 8 draws
Oracle 97.0%
Zero-shot CLM
Confidence
random
λ=0
Two-draft agreement
λ=0
Predeclared learner
Scored accuracy (%)
94
92
90
88
0
500
1000
1500
B Two-solver panel, Qwen3-0.6B
64 questions × 2 draws
Oracle 86.7%
λ=0
Confidence
random
80
75
70
65
60
0
500
1000
1500
C Two-solver panel, Qwen3-1.7B
64 questions × 2 draws
Oracle 89.8%
random
λ=0
Confidence
88
87
86
85
84
83
82
0
500
1000
1500
Output tokens per state (draft included)
Refitted at each λ from 0 to 0.3
Random escalation at equal cost
λ = 0.05, the primary cost
Always direct (left), always think (right)
Other rules in Table 1 at λ = 0.05
Hindsight oracle, off scale